Concurrency Control in Distributed Systems

Concurrency in Non-Distributed Systems:

Scenario:

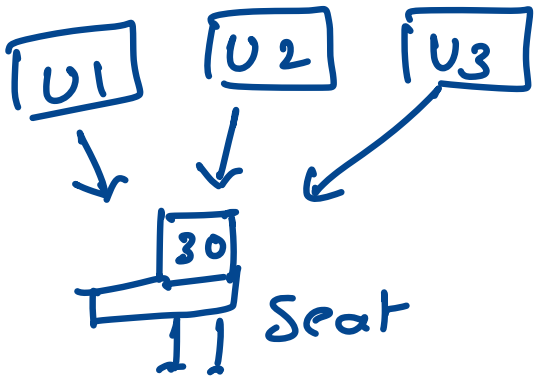

In Non-distributed systems, if we introduce synchronized block concept -- we may control concurrency issues.

Process

- Thread1
- Thread2
- Thread3
- ⋮

```
Synchronized() {
    if Free:
        Book seat
}
```

U1 → free — update to Booked

U2 → Not free — Nothing happen

U3 → Not free — ..

Concurrency in Distributed Systems:

In distributed systems, multiple processes run on multiple servers.

Now imagine multiple users accessing multiple processes to book a same seat. Here, synchronized block of code don't support. Therefore, we need to come up with a better strategy.

Before deep diving, we have to understand a set of concepts.

1) Transaction
2) DB Locking
3) Isolation levels.

Transaction: Set of instructions combined called a transaction.

If any instruction fails, the whole transaction has to revert back or roll back

T1 2

BankA -20 = BANKA

Failure X Bank B = B+20

Roll Back Commit

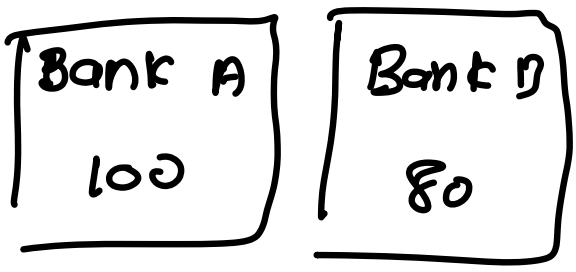

If failure encountered before commit then we need to roll back.

Here we have 2 consistent states

(1)

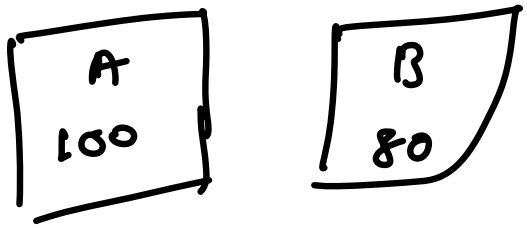

(2)

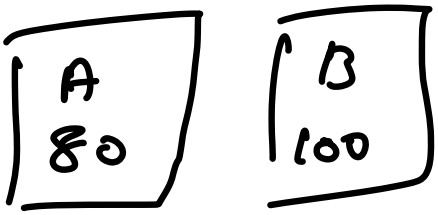

Remaining all are inconsistent states

DB Locking:

1) Shared Locking
2) Exclusive Locking.

Shared Locking: (Read Lock) (S)

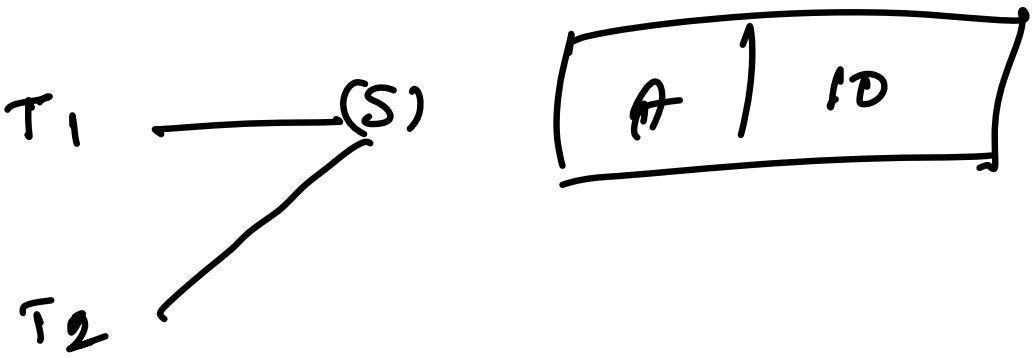

T1 & T2 can acquire shared locked and read data simultaneously.

Exclusive Lock: (Write lock) (X)

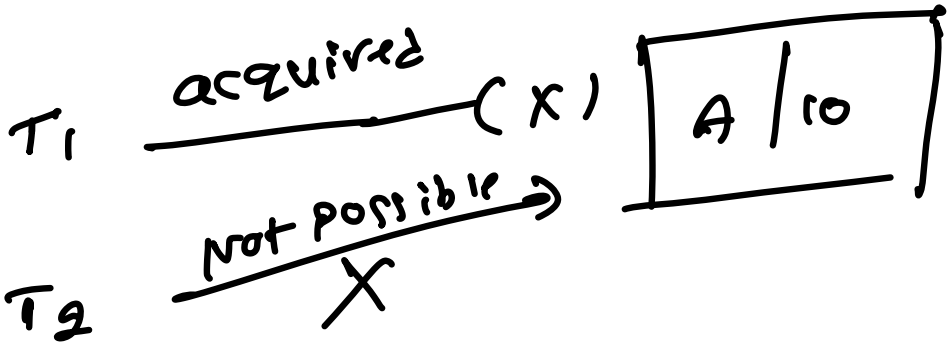

Write lock can be acquired one at a time

	has shared lock	has Exclusive lock
Shared Lock	✓	X
Exclusive Lock	X	X

1) if T_i is accessing shared lock then $T_2 \ldots T_n$ can access shared lock to read data
2) if T_i is accessing shared lock then $T_2 \ldots T_n$ cannot acquire exclusive lock.
3) If T_i is having exclusive lock then $T_2 \ldots T_n$ cannot acquire shared lock
4) If T_i is having exclusive lock then $T_2 \ldots T_n$ cannot acquire exclusive lock

Isolation levels in Distributed systems:

Before deep diving we need to understand

1) Dirty Read
2) Non-repeatable Read
3) Phantom read

Dirty Read:

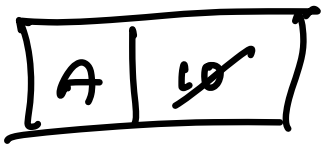

T1 {
Read A // 10
...
Read A // 11
...
commit
}

T2 {
Incrite A // 11
... X failure.
commit ;
}

Before committing write on DB (A) — T1 read the data updated by T2.

After Roll over -- 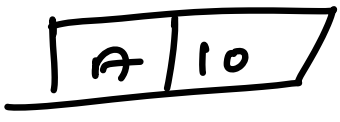

A becomes 10

but T1 did some calculations with A = 11

This is Dirty Read.

Non-Repeatable Read:

Read different values which are Non-Repeatable Committed

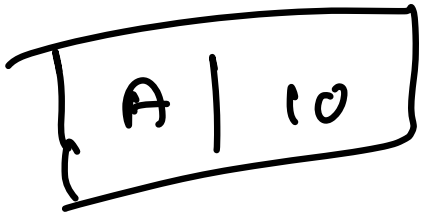

T1	T2
read A // 10	write A = 11
	Commit
read A // 11	

We read different values of A within a transaction

Phantom Read:

Read different no. of rows when a query is executed (committed)

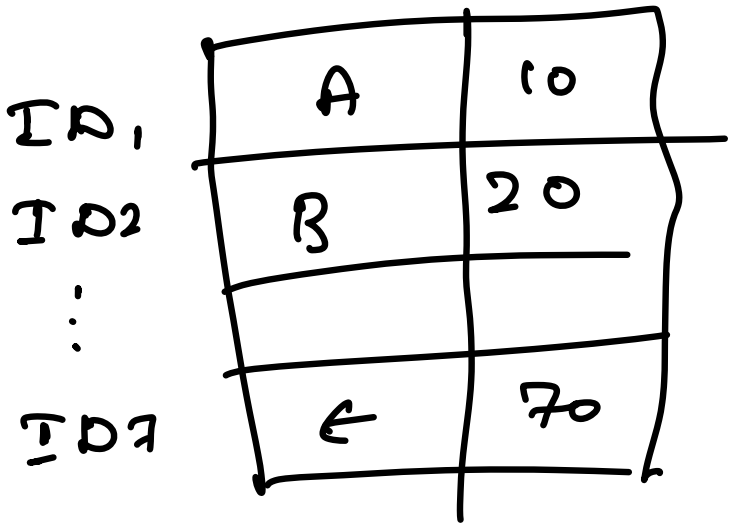

Select ID where ID>2 & ID≤5

First time read

We get

ID3, ID4, ID5

Now, some updates happened in DB.

We added a new row

Same query

select ID where ID>2 & ID≤5

we get

ID3, ID4, IDx, IDy, ID5

Types of isolation levels:

1) Read uncommited
2) Read committed
3) Repeatable Read
4) Serializable

Read Uncommitted:

As name suggest we need to support for reading uncommitted data.

So, We can support

1) Dirty Read
2) Non-repeatable Read
3) phantom Read

Isolation level	Dirty Read	Non-Repeatable read	Phantom Read
Read uncommitted	✓	✓	✓
Read Committed	X	✓	✓
Repeatable Read	X	X	✓
Serializable	X	X	X

Isolation level	Shared lock	Exclusive lock
Read uncommitted	X	X
Read Committed	Shared lock but release after read	Exclusive lock until write is done in a transaction
Repeatable Read	Shared lock till Transaction ends	Exclusive lock until Transaction ends
Serializable	Shared cor + range lock till Transaction ends	Range lock till Transaction ends

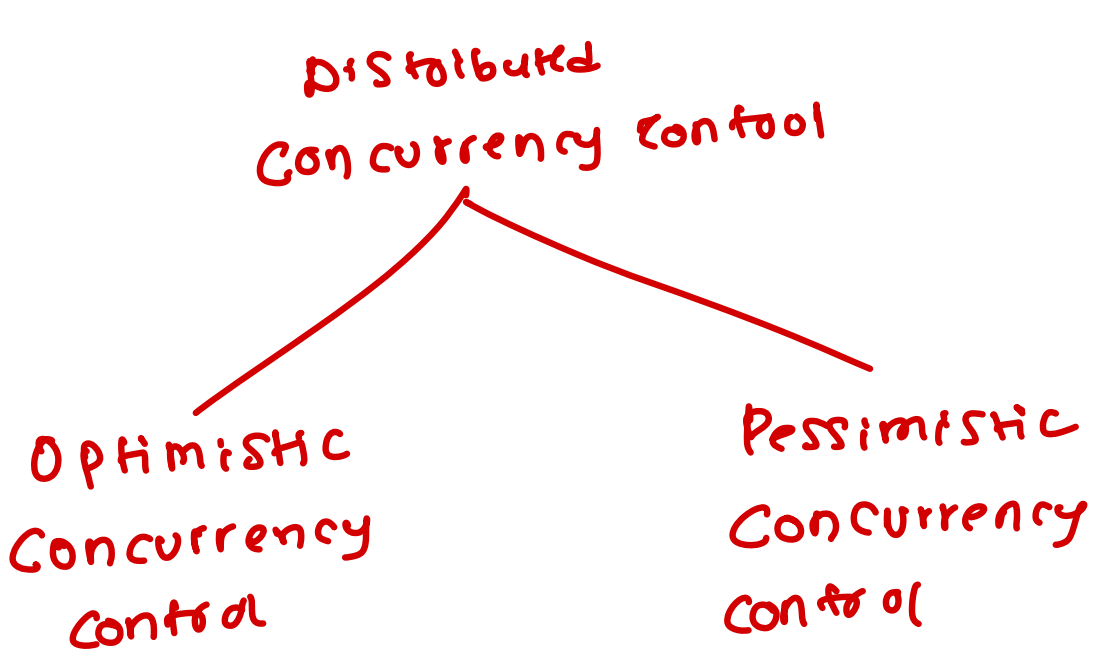

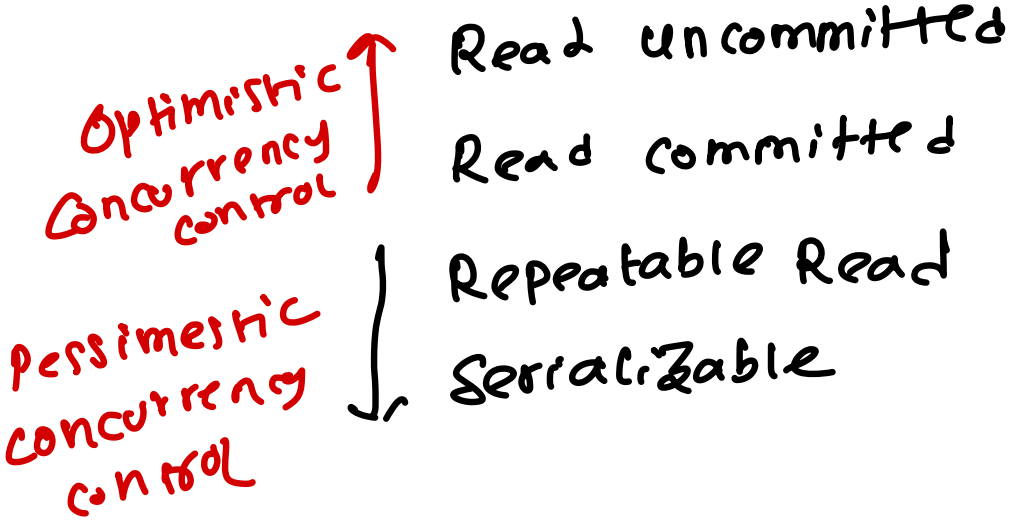

Choose distributed concurrency control as per the requirement in system design.